

```
In [26]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [27]: df=pd.read_csv(r"C:\Users\rubin\Downloads\fiat500_VehicleSelection_Dataset.xls.csv")
df
```

Out[27]:

|      | #NAME? | model  | engine_power | age_in_days | km     | previous_owners | lat       | lon       | price |
|------|--------|--------|--------------|-------------|--------|-----------------|-----------|-----------|-------|
| 0    | 1      | lounge | 51           | 882         | 25000  | 1               | 44.907242 | 8.611560  | 8900  |
| 1    | 2      | pop    | 51           | 1186        | 32500  | 1               | 45.666359 | 12.241890 | 8800  |
| 2    | 3      | sport  | 74           | 4658        | 142228 | 1               | 45.503300 | 11.417840 | 4200  |
| 3    | 4      | lounge | 51           | 2739        | 160000 | 1               | 40.633171 | 17.634609 | 6000  |
| 4    | 5      | pop    | 73           | 3074        | 106880 | 1               | 41.903221 | 12.495650 | 5700  |
| ...  | ...    | ...    | ...          | ...         | ...    | ...             | ...       | ...       | ...   |
| 1533 | 1534   | sport  | 51           | 3712        | 115280 | 1               | 45.069679 | 7.704920  | 5200  |
| 1534 | 1535   | lounge | 74           | 3835        | 112000 | 1               | 45.845692 | 8.666870  | 4600  |
| 1535 | 1536   | pop    | 51           | 2223        | 60457  | 1               | 45.481541 | 9.413480  | 7500  |
| 1536 | 1537   | lounge | 51           | 2557        | 80750  | 1               | 45.000702 | 7.682270  | 5990  |
| 1537 | 1538   | pop    | 51           | 1766        | 54276  | 1               | 40.323410 | 17.568270 | 7900  |

1538 rows × 9 columns

```
In [28]: df.head()
```

Out[28]:

|   | #NAME? | model  | engine_power | age_in_days | km     | previous_owners | lat       | lon       | price |
|---|--------|--------|--------------|-------------|--------|-----------------|-----------|-----------|-------|
| 0 | 1      | lounge | 51           | 882         | 25000  | 1               | 44.907242 | 8.611560  | 8900  |
| 1 | 2      | pop    | 51           | 1186        | 32500  | 1               | 45.666359 | 12.241890 | 8800  |
| 2 | 3      | sport  | 74           | 4658        | 142228 | 1               | 45.503300 | 11.417840 | 4200  |
| 3 | 4      | lounge | 51           | 2739        | 160000 | 1               | 40.633171 | 17.634609 | 6000  |
| 4 | 5      | pop    | 73           | 3074        | 106880 | 1               | 41.903221 | 12.495650 | 5700  |

```
In [29]: df.shape
```

Out[29]: (1538, 9)

In [30]: df.describe

```
Out[30]: <bound method NDFrame.describe of
0      1 lounge      51      882  25000      1 \
1      2 pop        51     1186  32500      1
2      3 sport      74     4658 142228      1
3      4 lounge     51     2739 160000      1
4      5 pop        73     3074 106880      1
...    ...    ...    ...    ...    ...
1533   1534 sport     51     3712 115280      1
1534   1535 lounge     74     3835 112000      1
1535   1536 pop        51     2223  60457      1
1536   1537 lounge     51     2557  80750      1
1537   1538 pop        51     1766  54276      1

      lat      lon  price
0    44.907242  8.611560  8900
1    45.666359 12.241890  8800
2    45.503300 11.417840  4200
3    40.633171 17.634609  6000
4    41.903221 12.495650  5700
...    ...    ...    ...
1533  45.069679  7.704920  5200
1534  45.845692  8.666870  4600
1535  45.481541  9.413480  7500
1536  45.000702  7.682270  5990
1537  40.323410 17.568270  7900

[1538 rows x 9 columns]>
```


In [31]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   #NAME?                 1538 non-null  int64
1   model                  1538 non-null  object
2   engine_power           1538 non-null  int64
3   age_in_days            1538 non-null  int64
4   km                     1538 non-null  int64
5   previous_owners        1538 non-null  int64
6   lat                   1538 non-null  float64
7   lon                   1538 non-null  float64
8   price                  1538 non-null  int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [32]: df.isna().any()

```
Out[32]: #NAME?      False
model          False
engine_power    False
age_in_days     False
km              False
previous_owners False
lat             False
lon             False
price           False
dtype: bool
```

```
In [33]: sns.lmplot(x='km',y='price',data=df,order=2,ci=None)
plt.show()
```



```
In [34]: x=np.array(df['km']).reshape(-1,1)
y=np.array(df['price']).reshape(-1,1)
```


```
In [35]: df.dropna(inplace=True)
```

```
In [36]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
#splitting data into train and test
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```


0.7002560851710697

```
In [37]: y_pred=regr.predict(x_test)

plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



```
In [38]: df500=df[:][:500]
sns.lmplot(x="km",y="price",data=df500,order=1,ci=None)
plt.show()
```



```
In [39]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

```
In [40]: #train model
model=LinearRegression()
model.fit(x_train,y_train)
#Evaluation the model on the test set
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.7002560851710697

```
In [41]: from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [42]: features=df.columns[0:1]
target=df.columns[-1]
```

```
In [43]: converter={"model":{"sport":1,"loung":2,"pop":3}}
df=df.replace(converter)
df
```


Out[43]:

|      | #NAME? | model | engine_power | age_in_days | km     | previous_owners | lat       | lon       | price |
|------|--------|-------|--------------|-------------|--------|-----------------|-----------|-----------|-------|
| 0    | 1      | 2     | 51           | 882         | 25000  | 1               | 44.907242 | 8.611560  | 8900  |
| 1    | 2      | 3     | 51           | 1186        | 32500  | 1               | 45.666359 | 12.241890 | 8800  |
| 2    | 3      | 1     | 74           | 4658        | 142228 | 1               | 45.503300 | 11.417840 | 4200  |
| 3    | 4      | 2     | 51           | 2739        | 160000 | 1               | 40.633171 | 17.634609 | 6000  |
| 4    | 5      | 3     | 73           | 3074        | 106880 | 1               | 41.903221 | 12.495650 | 5700  |
| ...  | ...    | ...   | ...          | ...         | ...    | ...             | ...       | ...       | ...   |
| 1533 | 1534   | 1     | 51           | 3712        | 115280 | 1               | 45.069679 | 7.704920  | 5200  |
| 1534 | 1535   | 2     | 74           | 3835        | 112000 | 1               | 45.845692 | 8.666870  | 4600  |
| 1535 | 1536   | 3     | 51           | 2223        | 60457  | 1               | 45.481541 | 9.413480  | 7500  |
| 1536 | 1537   | 2     | 51           | 2557        | 80750  | 1               | 45.000702 | 7.682270  | 5990  |
| 1537 | 1538   | 3     | 51           | 1766        | 54276  | 1               | 40.323410 | 17.568270 | 7900  |

1538 rows × 9 columns

```
In [46]: plt.figure(figsize = (10, 10))
sns.heatmap(df.corr(), annot = True)
```

Out[46]: <Axes: >



```
In [47]: #X and y values
X = df[features].values
y = df[target].values
#Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X\_train is (1153, 1)  
The dimension of X\_test is (385, 1)

```
In [48]: lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.00310286926477088

The test score for lr model is -0.008405634316406507

```
In [49]: #Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test scorefor ridge regression
train_score_ridge = ridgeReg.score(X_train, y_train)
test_score_ridge = ridgeReg.score(X_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.0031026398591535997

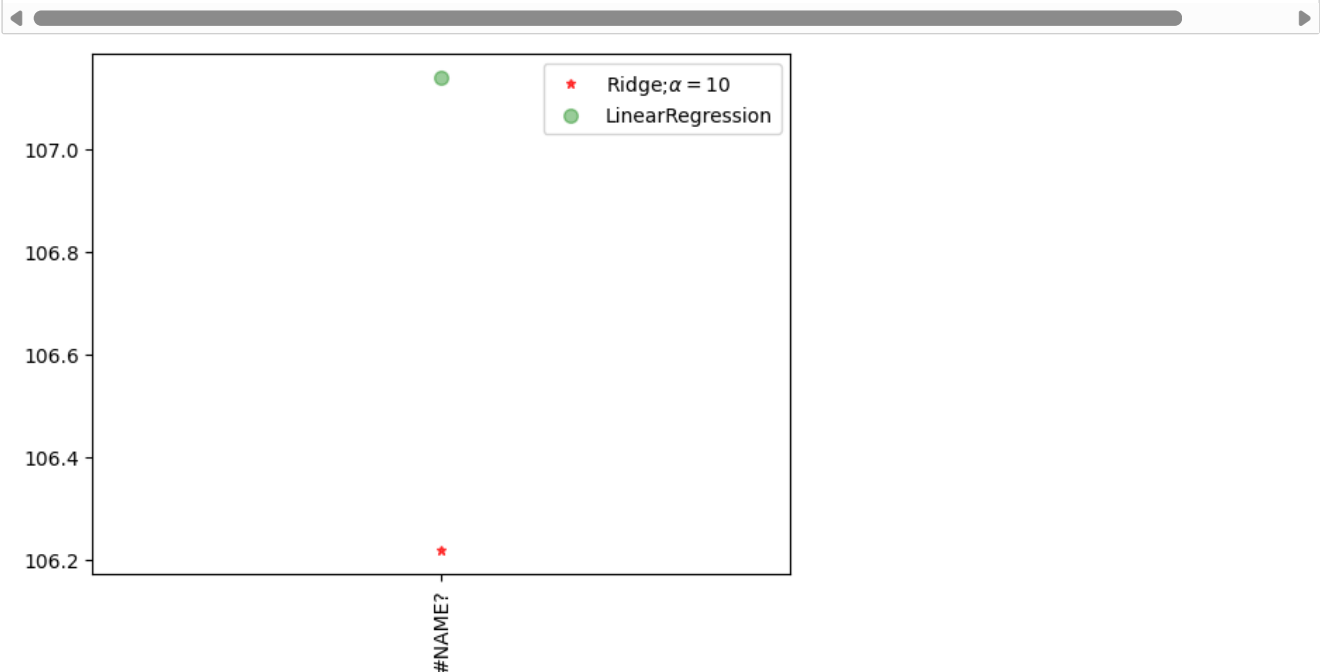
The test score for ridge model is -0.008307809466001403

```
In [50]: plt.figure(figsize=(10,10))
```

Out[50]: <Figure size 1000x1000 with 0 Axes>

<Figure size 1000x1000 with 0 Axes>

```
In [52]: plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color='red',label=r'Ridge;\alpha=10$')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker="o",markersize=7,color='green',label='LinearRegression')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```



```
In [53]: #Lasso regression model
print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.003075838461310987

The test score for ls model is -0.007367578602064606

```
In [54]: pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[54]: <Axes: >



```
In [55]: #Using the linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
#score
print(lasso_cv.score(X_train, y_train))
print(lasso_cv.score(X_test, y_test))
```

0.00310259895673648

-0.008299466692577973




In [56]: #plot size

```

plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge; $\alpha=10$')
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso; $\alpha = \text{grid}$')
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()

```



In [57]: #Using the Linear CV model

```

from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))

```

The train score for ridge model is 0.0031026398591535997  
 The train score for ridge model is -0.008307809466002958

In [ ]: